



# Scalable data processing on computer clusters

**Gergely Lukács**

Pázmány Péter Catholic University

Faculty of Information Technology



Budapest, Hungary

lukacs@itk.ppke.hu

# Problem with amount of data?

- Better algorithms
- Numba, C/Cython
- More efficient formats for data
- Data subsampling
  
- Laptop (4 cores, 16GB RAM, 1 TB hdd)
- Server (24 cores, 1TB RAM)
- Parallelization
  - Methods/architectures for parallelization
- HW accelerators (depending on task: GPGPU, FPGA)

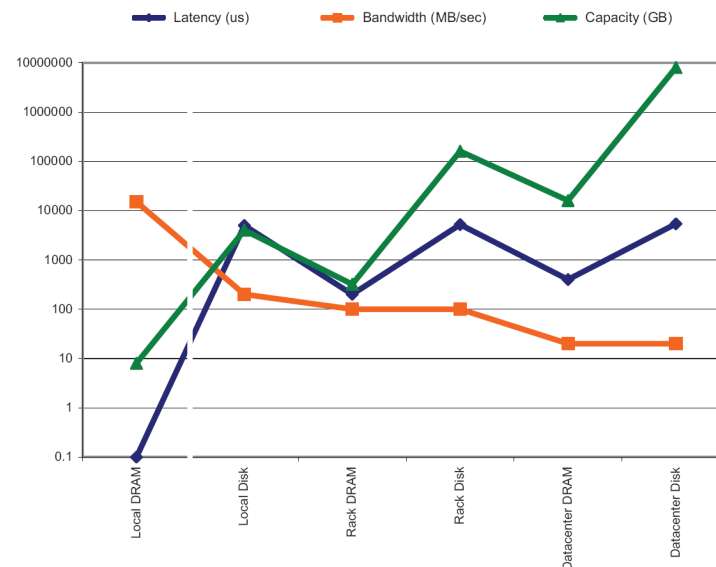

# Building Blocks



# Major ideas

- Process data sequentially, avoid random access
  - Seeks are expensive, disk throughput is reasonable
- Move processing to the data
  - Cluster have limited bandwidth
- Seamless scalability
- From the mythical man-month to the tradable machine-hour
  - Automatic parallelization & distribution
  - Fault tolerance
  - I/O scheduling
  - Monitoring & status updates

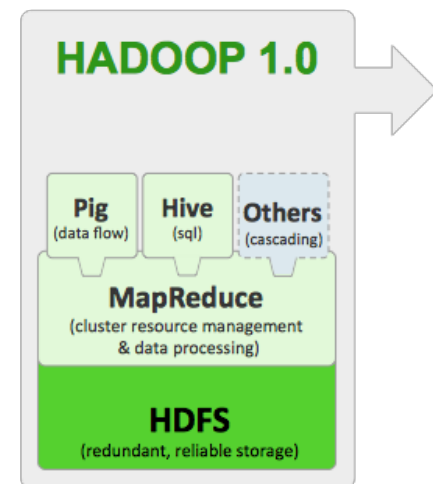
# Storage Hierarchy, Seek vs. scan



- Consider a 1 TB database with 100 byte records
  - We want to update 1 percent of the records
- Scenario 1: random access
  - Each update takes ~30 ms (seek, read, write)
  - $10^8$  updates = ~35 days
- Scenario 2: rewrite all records
  - Assume 100 MB/s throughput
  - Time = 5.6 hours(!)
- Lesson: avoid random seeks!

# Hadoop (1.0), MapReduce

- Elosztott fájlrendszer: HDFS
  - Skálázódás, hibatűrés
  - Elosztott, redundáns
- MapReduce



# Map/Reduce

- Map/Reduce
  - Programming model from Lisp  
(and other functional languages)
- Many problems can be phrased this way
- Easy to distribute across nodes
- Nice retry/failure semantics

# Map in Lisp (Scheme)

- (map *f list* [*list<sub>2</sub> list<sub>3</sub> ...*])

Unary operator

- (map square '(1 2 3 4))

- (1 4 9 16)

Binary operator

- (reduce + '(1 4 9 16))

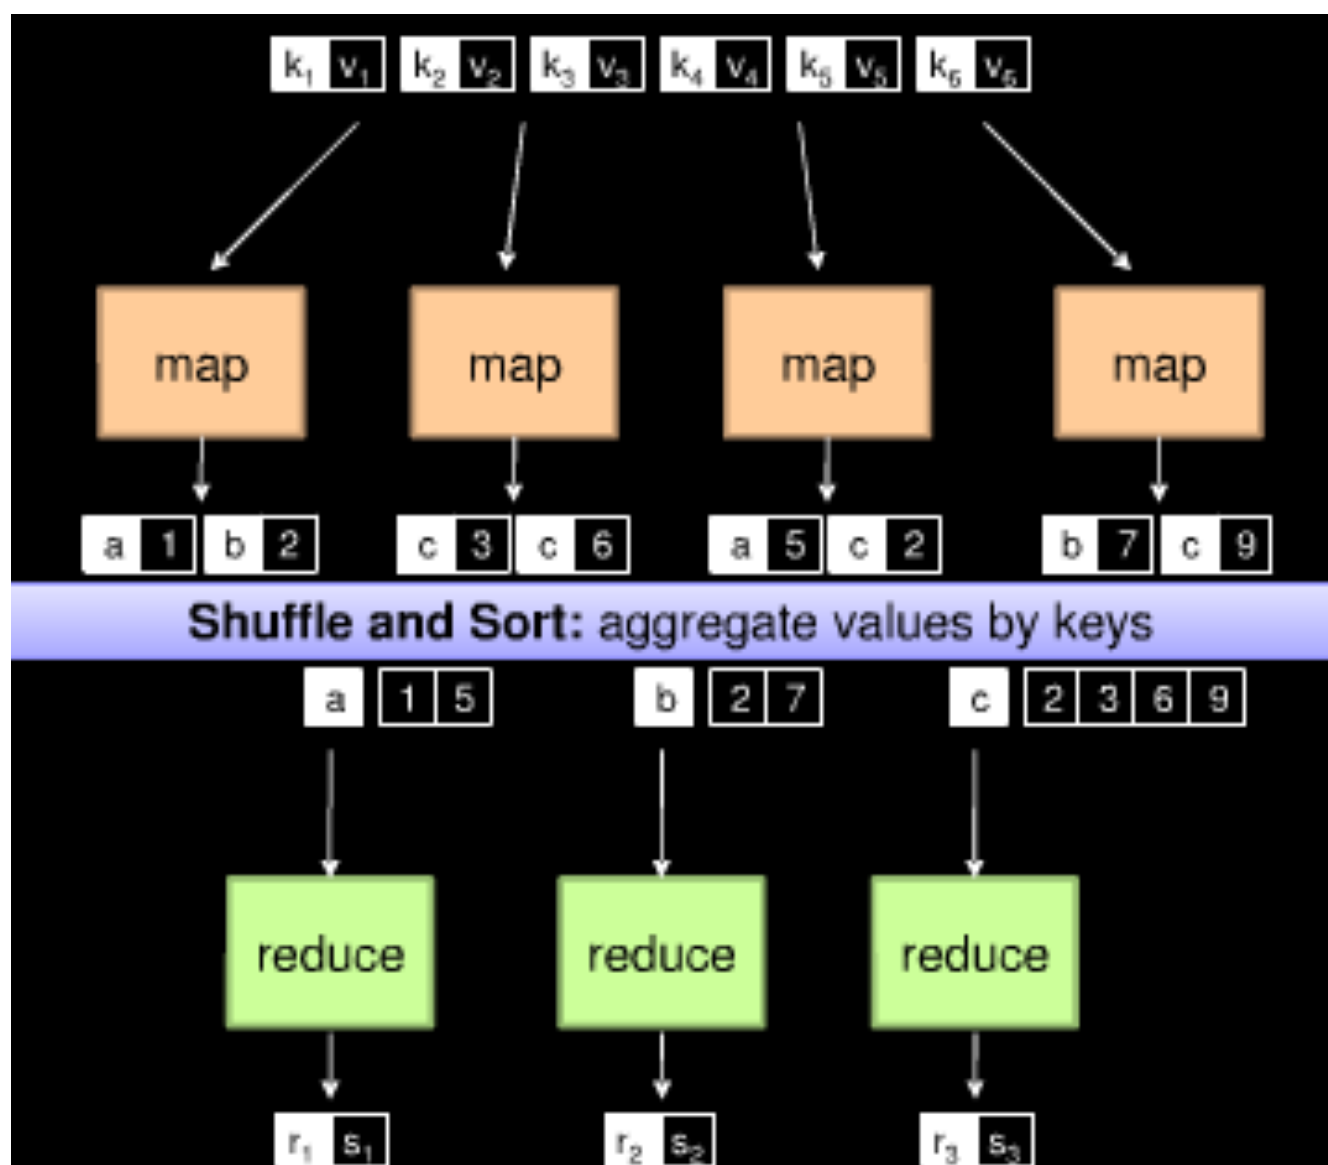
- 30

- (reduce + (map square (map - l<sub>1</sub> l<sub>2</sub>))))



# Map/Reduce ala Google

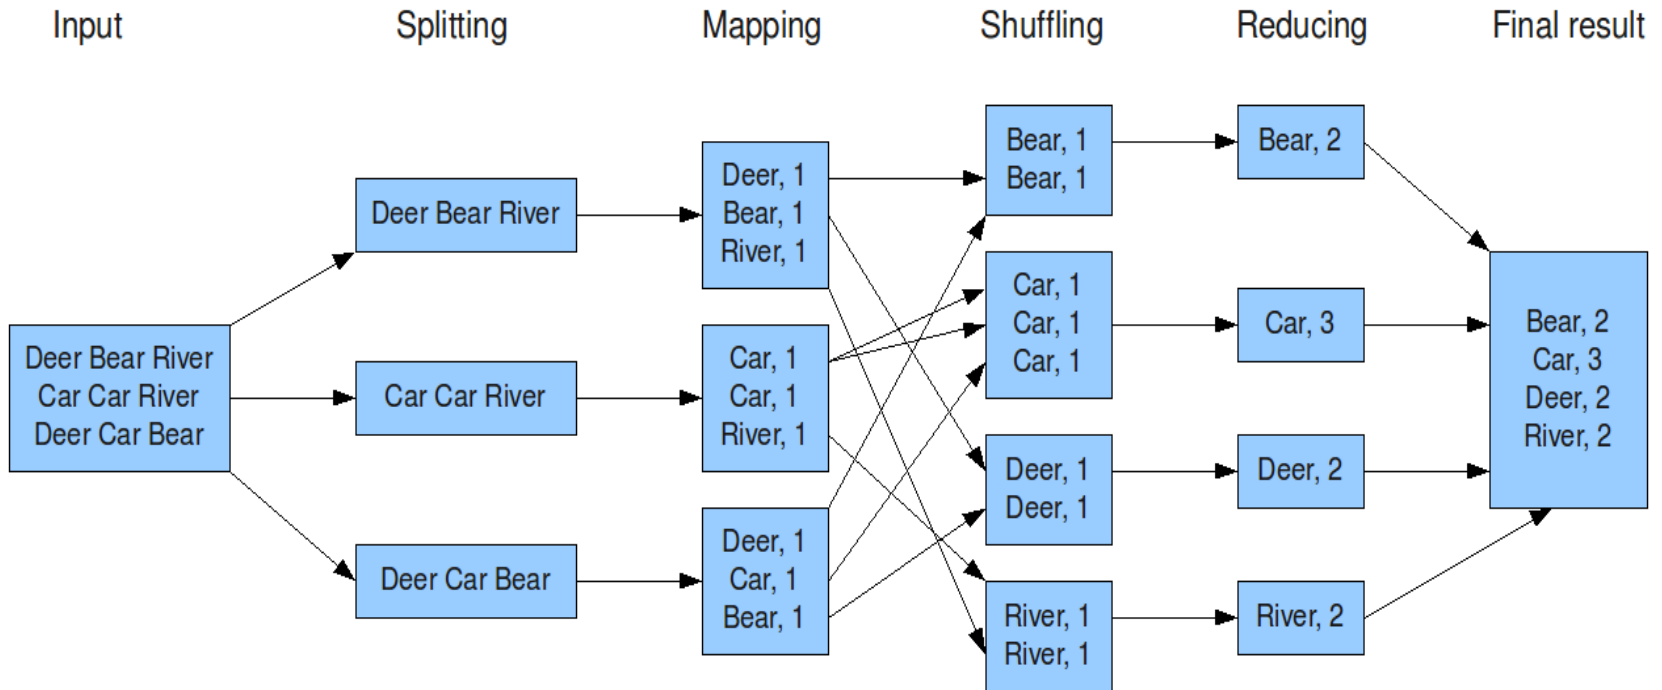
- `map(key, val)` is run on each item in set
  - emits new-key / new-val pairs
- `reduce(key, vals)`
  - All values with the same key are reduced together
  - emits final output



# count words in docs


- Input consists of (url, contents) pairs
- map(key=url, val=contents):
  - For each word  $w$  in contents, emit ( $w$ , “1”)
- reduce(key=word, values=uniq\_counts):
  - Sum all “1”s in values list
  - Emit result “(word, sum)”

# The overall MapReduce word count process



# Model is Widely Applicable

## MapReduce Programs In Google Source Tree



### Example uses:

distributed grep

term-vector / host

document clustering

...

distributed sort

web access log stats

machine learning

...

web link-graph reversal

inverted index construction

statistical machine translation

...

# Execution Framework

- The execution framework handles everything else...
  - Scheduling: assigns workers to map and reduce tasks
  - “Data distribution”: moves processes to data
  - Synchronization: gathers, sorts, and shuffles intermediate data
  - Errors and faults: detects worker failures and restarts
- Limited control over data and execution flow
  - All algorithms must be expressed in m, r, c, p
- You don’t know:
  - Where mappers and reducers run
  - When a mapper or reducer begins or finishes
  - Which input a particular mapper is processing
  - Which intermediate key a particular reducer is processing

# In MapReduce

```

import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.Writable;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.RecordReader;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String key = line.substring(0, firstComma);
            String value = line.substring(firstComma + 1);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text(key + value);
            oc.collect(outKey, outVal);
        }

        public static class LoadAndFilterUsers extends MapReduceBase
            implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(1);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2" + value);
            oc.collect(outKey, outVal);
        }

        public static class Join extends MapReduceBase
            implements Reducer<Text, Text, Text, Text> {

        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();

            while (iter.hasNext()) {
                Text t = iter.next();
                String val = t.toString();
                if (val.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }

            reporter.setStatus("OK");
        }

        // Do the cross product and collect the values
        for (String s1 : first) {
            for (String s2 : second) {
                String outval = key + "," + s1 + "," + s2;
                oc.collect(null, new Text(outval));
                reporter.setStatus("OK");
            }
        }
    }

    public static class LoadJoined extends MapReduceBase
        implements Mapper<Text, Text, Text, LongWritable> {

        public void map(
            Text k,
            Text val,
            OutputCollector<Text, LongWritable> oc,
            Reporter reporter) throws IOException {
            // Find the url
            String line = val.toString();
            int firstComma = line.indexOf(',');
            int secondComma = line.indexOf(',', firstComma + 1);
            String key = line.substring(firstComma, secondComma);
            // drop the rest of the record, I don't need it anymore
            // just pass a 1 for the combiner/reducer to sum
            Text outKey = new Text(key);
            Text outVal = new LongWritable(1L);
            oc.collect(outKey, outVal);
        }

        public static class ReduceUrls extends MapReduceBase
            implements Reducer<Text, LongWritable, WritableComparable, Writable> {

        public void reduce(
            Text key,
            Iterator<LongWritable> iter,
            OutputCollector<WritableComparable, Writable> oc,
            Reporter reporter) throws IOException {
            // Add up all the values we see
            long sum = 0;
            while (iter.hasNext()) {
                sum += iter.next().get();
            }
            reporter.setStatus("OK");
            oc.collect(key, new LongWritable(sum));
        }

        public static class LoadClicks extends MapReduceBase
            implements Mapper<WritableComparable, Writable, LongWritable, Text> {

        public void map(
            WritableComparable key,
            Writable val,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            oc.collect((LongWritable)val, (Text)key);
        }

        public static class LimitClicks extends MapReduceBase
            implements Reducer<LongWritable, Text, LongWritable, Text> {

        int count = 0;

        public void reduce(
            LongWritable key,
            Iterator<Text> iter,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            // Only output the first 100 records
            while (count < 100 && iter.hasNext()) {
                oc.collect(key, iter.next());
                count++;
            }
        }

        public static void main(String[] args) throws IOException {
            JobConf lp = new JobConf(MRExample.class);
            lp.setJobName("Load Pages");
            lp.setInputFormat(TextInputFormat.class);
            lp.setOutputFormat(TextOutputFormat.class);
            lp.setMapperClass(LoadPages.class);
            FileInputFormat.addInputPath(lp, new Path("/user/gates/pages"));
            FileOutputFormat.setOutputPath(lp, new Path("/user/gates/tmp/indexed_pages"));
            lp.setNumReduceTasks(0);
            Job loadPages = new Job(lp);

            JobConf lfu = new JobConf(MRExample.class);
            lfu.setJobName("Load and Filter Users");
            lfu.setInputFormat(SequenceFileInputFormat.class);
            lfu.setOutputFormat(SequenceFileOutputFormat.class);
            lfu.setMapperClass(LoadAndFilterUsers.class);
            lfu.setReducerClass(Join.class);
            FileInputFormat.addInputPath(lfu, new Path("/user/gates/users"));
            FileOutputFormat.setOutputPath(lfu, new Path("/user/gates/tmp/filtered_users"));
            lfu.setNumReduceTasks(0);
            Job loadUsers = new Job(lfu);

            JobConf join = new JobConf(MRExample.class);
            join.setJobName("Join Users and Pages");
            join.setInputFormat(SequenceFileInputFormat.class);
            join.setMapperClass(LoadPages.class);
            join.setReducerClass(Join.class);
            FileInputFormat.addInputPath(join, new Path("/user/gates/tmp/indexed_pages"));
            FileInputFormat.addInputPath(join, new Path("/user/gates/tmp/filtered_users"));
            FileOutputFormat.setOutputPath(join, new Path("/user/gates/tmp/joined"));
            join.setNumReduceTasks(50);
            Job joinJob = new Job(join);
            joinJob.addDependingJob(loadPages);
            joinJob.addDependingJob(loadUsers);

            JobConf group = new JobConf(MRExample.class);
            group.setJobName("Group URLs");
            group.setInputFormat(SequenceFileInputFormat.class);
            group.setOutputFormat(SequenceFileOutputFormat.class);
            group.setMapperClass(LoadJoined.class);
            group.setCombinerClass(Join.class);
            group.setReducerClass(ReduceUrls.class);
            FileInputFormat.addInputPath(group, new Path("/user/gates/tmp/joined"));
            FileOutputFormat.setOutputPath(group, new Path("/user/gates/tmp/grouped"));
            group.setNumReduceTasks(50);
            Job groupJob = new Job(group);
            groupJob.addDependingJob(joinJob);

            JobConf top100 = new JobConf(MRExample.class);
            top100.setJobName("Top 100 sites");
            top100.setInputFormat(SequenceFileInputFormat.class);
            top100.setOutputFormat(SequenceFileOutputFormat.class);
            top100.setMapperClass(LoadClicks.class);
            top100.setCombinerClass(LimitClicks.class);
            top100.setReducerClass(LimitClicks.class);
            FileInputFormat.addInputPath(top100, new Path("/user/gates/tmp/grouped"));
            FileOutputFormat.setOutputPath(top100, new Path("/user/gates/top100/sitesForUsers18to25"));
            top100.setNumReduceTasks(1);
            Job limit = new Job(top100);
            limit.addDependingJob(groupJob);

            JobControl jc = new JobControl("First Top 100");
            jc.addJob(loadPages);
            jc.addJob(loadUsers);
            jc.addJob(joinJob);
            jc.addJob(groupJob);
            jc.addJob(limit);
            jc.run();
        }
    }
}

```

# In Pig Latin

```
Users      = load 'users' as (name, age);
Filtered   = filter Users by
              age >= 18 and age <= 25;
Pages      = load 'pages' as (user, url);
Joined     = join Filtered by name, Pages by user;
Grouped    = group Joined by url;
Summed     = foreach Grouped generate group,
              count(Joined) as clicks;
Sorted     = order Summed by clicks desc;
Top5       = limit Sorted 5;

store Top5 into 'top5sites';
```



# Writing an Hadoop MapReduce Program in Python

In this tutorial I will describe how to write a simple [MapReduce](#) program for [Hadoop](#) in the [Python](#) programming language.

## Motivation

Even though the Hadoop framework is written in Java, programs for Hadoop need not to be coded in Java but can also be developed in other languages like Python or C++ (the latter since version 0.14.1). However, [Hadoop's documentation](#) and the most prominent [Python example](#) on the Hadoop website could make you think that you *must* translate your Python code using [lython](#) into a Java jar file. Obviously, this is not very convenient and can even

### Table of Contents

- [Motivation](#)
- [What we want to do](#)
- [Prerequisites](#)
- [Python MapReduce Code](#)
  - [Map step: mapper.py](#)
  - [Reduce step: reducer.py](#)
  - [Test your code \(cat data | map | sort | reduce\)](#)
- [Running the Python Code on Hadoop](#)
  - [Download example input data](#)
  - [Copy local example data to HDFS](#)
  - [Run the MapReduce job](#)
- [Improved Mapper and Reducer code: using Python iterators and generators](#)
  - [mapper.py](#)
  - [reducer.py](#)

### About Me



I am a software engineer turned product manager based in Switzerland, Europe. In my day job I am working on products at [Confluent](#), the US startup founded by the creators of [Apache Kafka](#). [Read more »](#)

### Contact

✉ [michael@michael-noll.com](mailto:michael@michael-noll.com)

### Follow Me

🐦 Twitter

📡 Blog RSS

🐙 GitHub

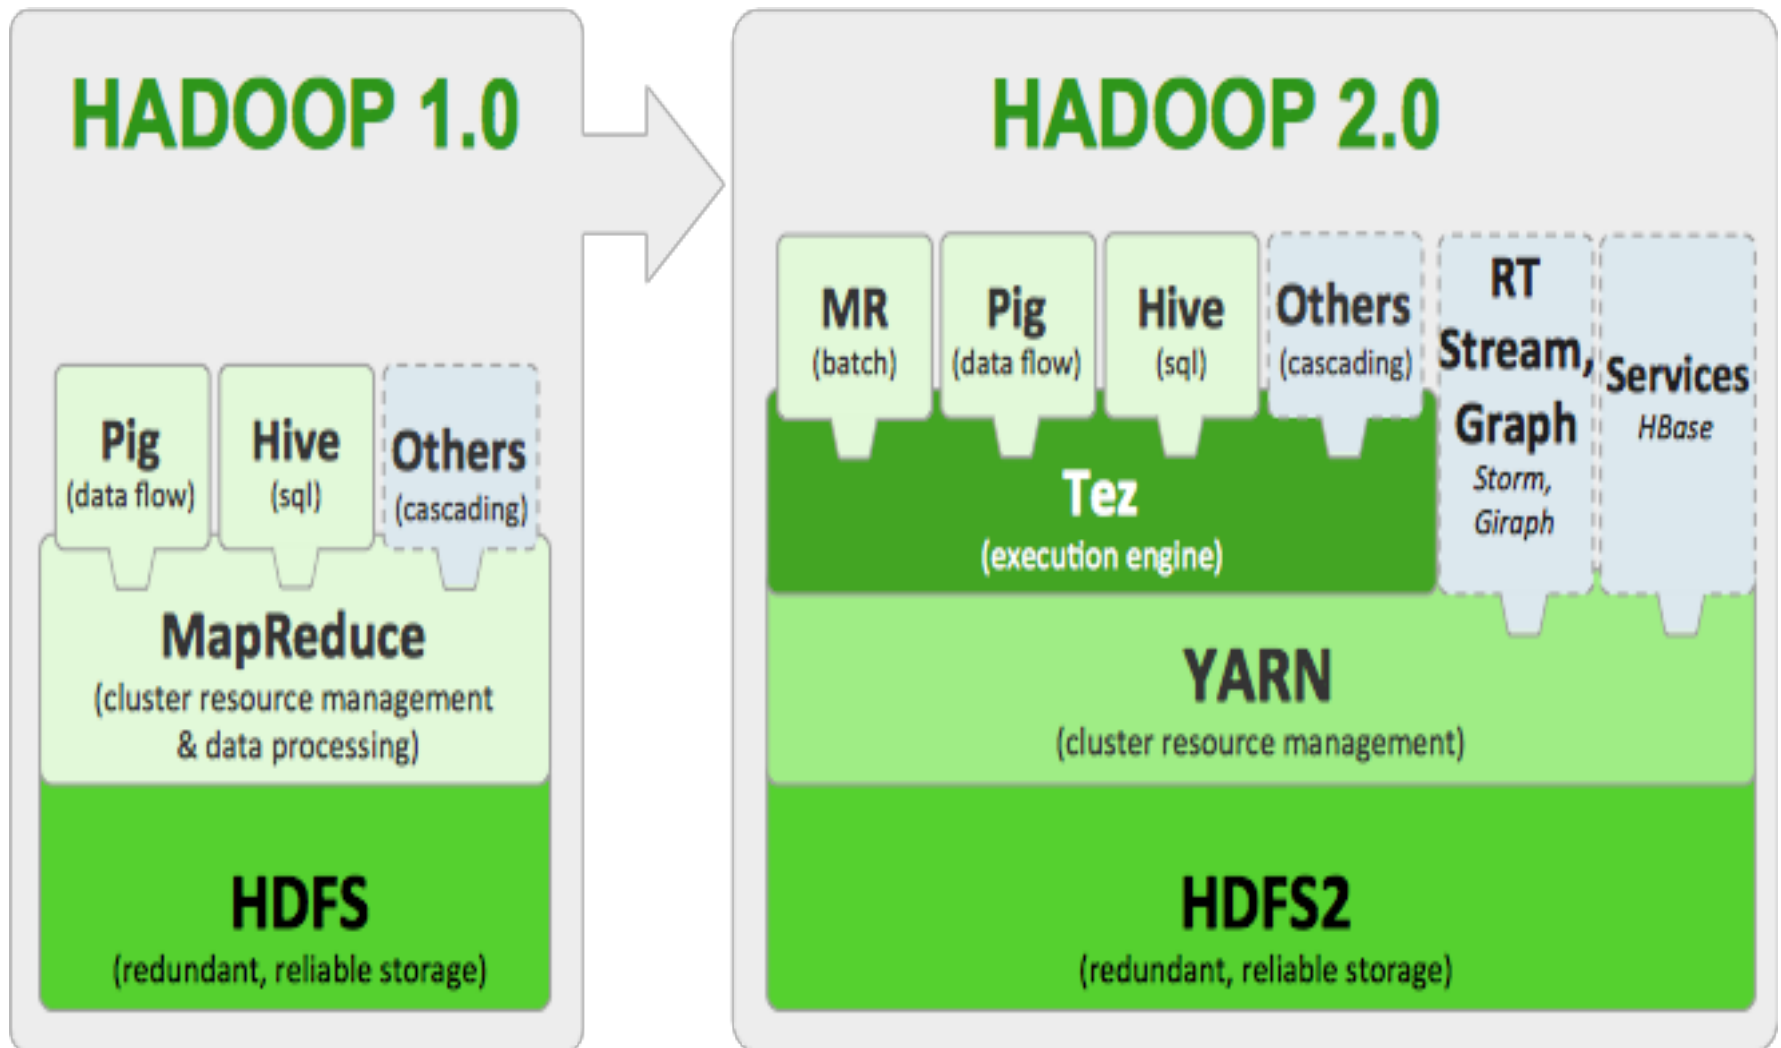
📺 Slideshare

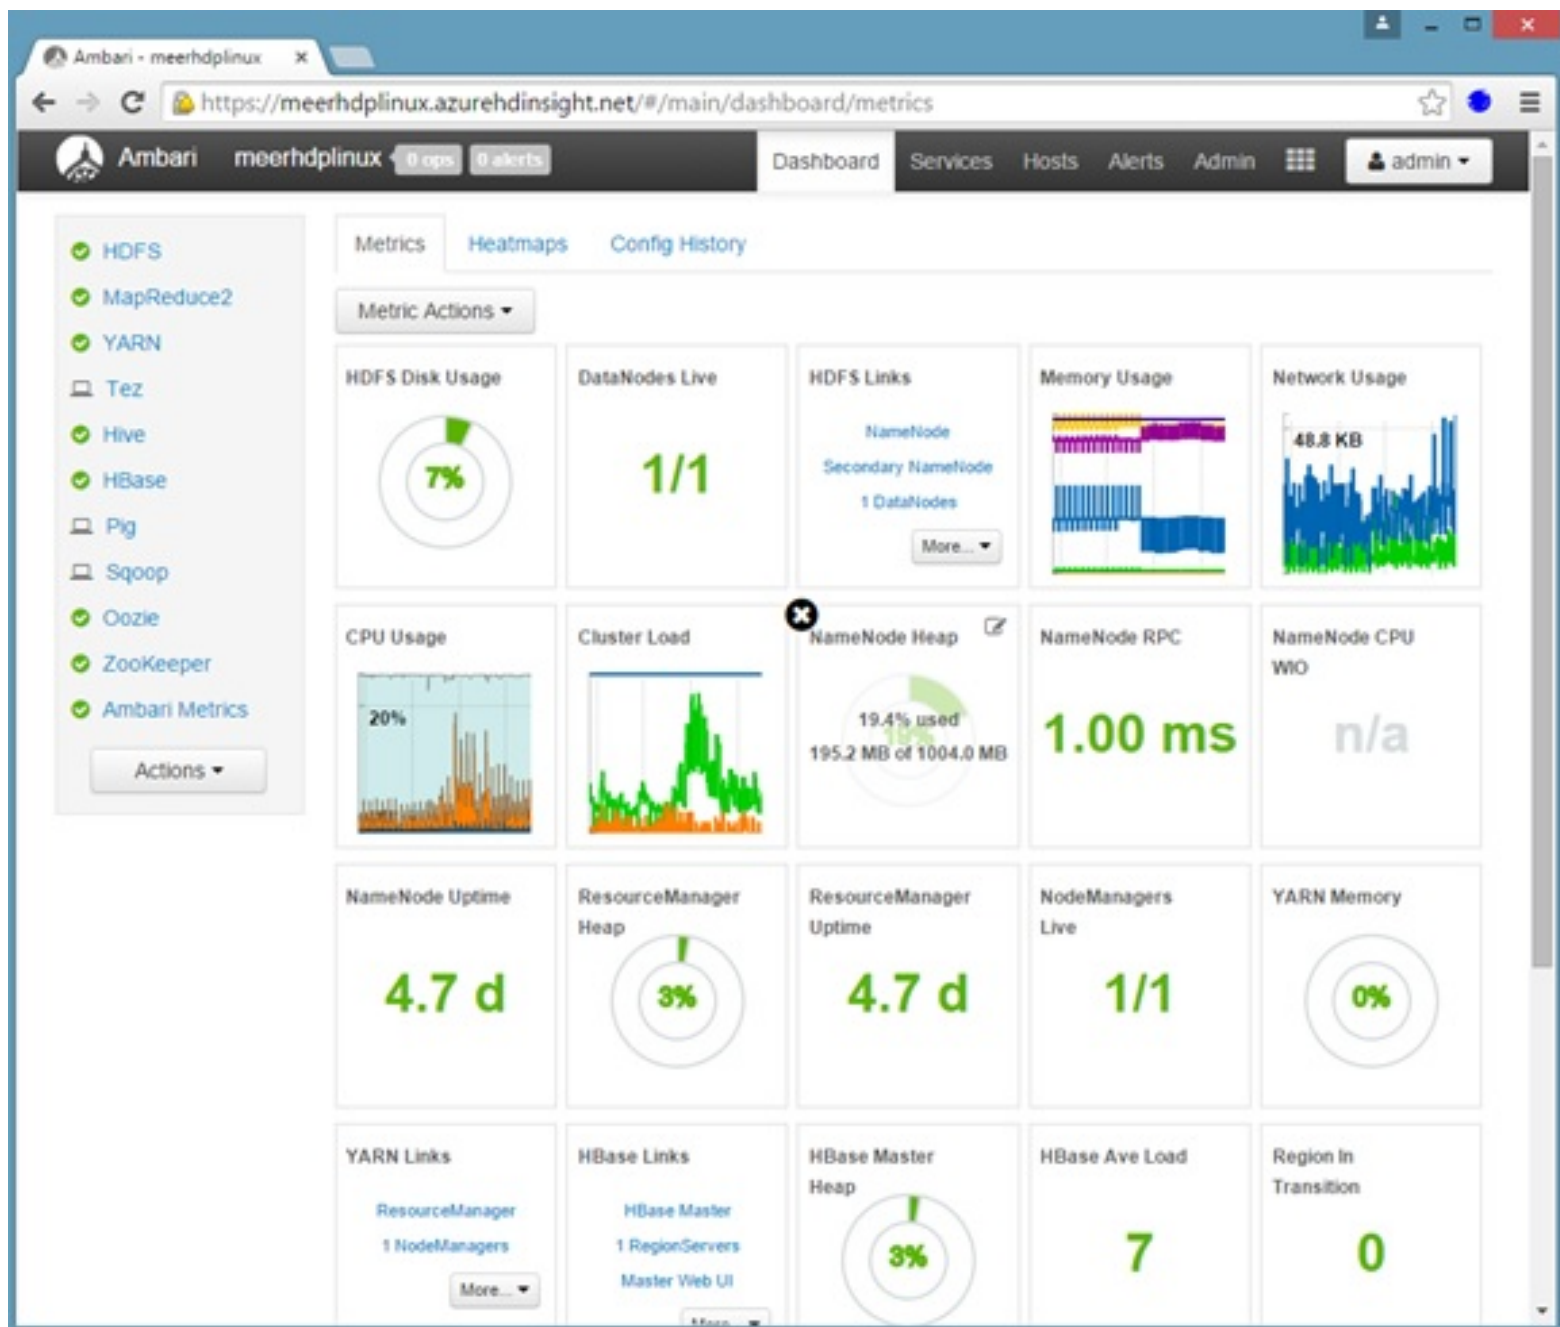
### Recent Posts

2010

# Limitations of MapReduce

- Linear dataflow:
  - read **input data from disk**,
  - map, reduce...,
  - **store result on disk**
- **Iterations** - when you need to process data again and again
- When your processing requires lot of data to be **shuffled** over the network.
- **Real-time** processing.
- Complex algorithms
- Processing graphs
- ....






Apache Spark™ - Lightning...

spark.apache.org

Search

☆📁🔔⬇️🏠😊Z📄🌐☰

GrapefonMost VisitedGetting StartedMNO - Címlapheadsets\_for\_rorhscha...szerverek



Spark™ Lightning-fast cluster computing

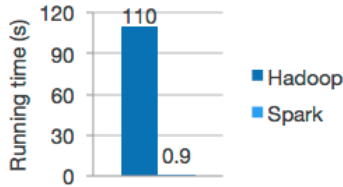
DownloadLibraries ▾Documentation ▾ExamplesCommunity ▾FAQApache Software Foundation ▾

Apache Spark™ is a fast and general engine for large-scale data processing.

Speed

Run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk.

Spark has an advanced DAG execution engine that supports cyclic data flow and in-memory computing.



Framework	Running time (s)
Hadoop	110
Spark	0.9

Logistic regression in Hadoop and Spark

Ease of Use

Write applications quickly in Java, Scala, Python, R.

```
text_file = spark.textFile("hdfs://...")
text_file.flatMap(lambda line: line.split())
            .map(lambda word: (word, 1))
```

Latest News

Spark Summit (June 6, 2016, San Francisco) agenda posted (Apr 17, 2016)

Spark 1.6.1 released (Mar 09, 2016)

Submission is open for Spark Summit San Francisco (Feb 11, 2016)

Spark Summit East (Feb 16, 2016, New York) agenda posted (Jan 14, 2016)

Archive

Download Spark

Built-in Libraries:

[SQL and DataFrames](#)


[Spark Streaming](#)

# Spark

- Resilient Distributed Dataset (RDD)
  - Read-Only
  - Distributed
  - Fault-tolerant
  - Caching can be controlled
  - Iterative algorithms
  - Interactive/exploratory data analysis

# Spark

- Standalone (native Spark cluster)
- Hadoop YARN
- ...



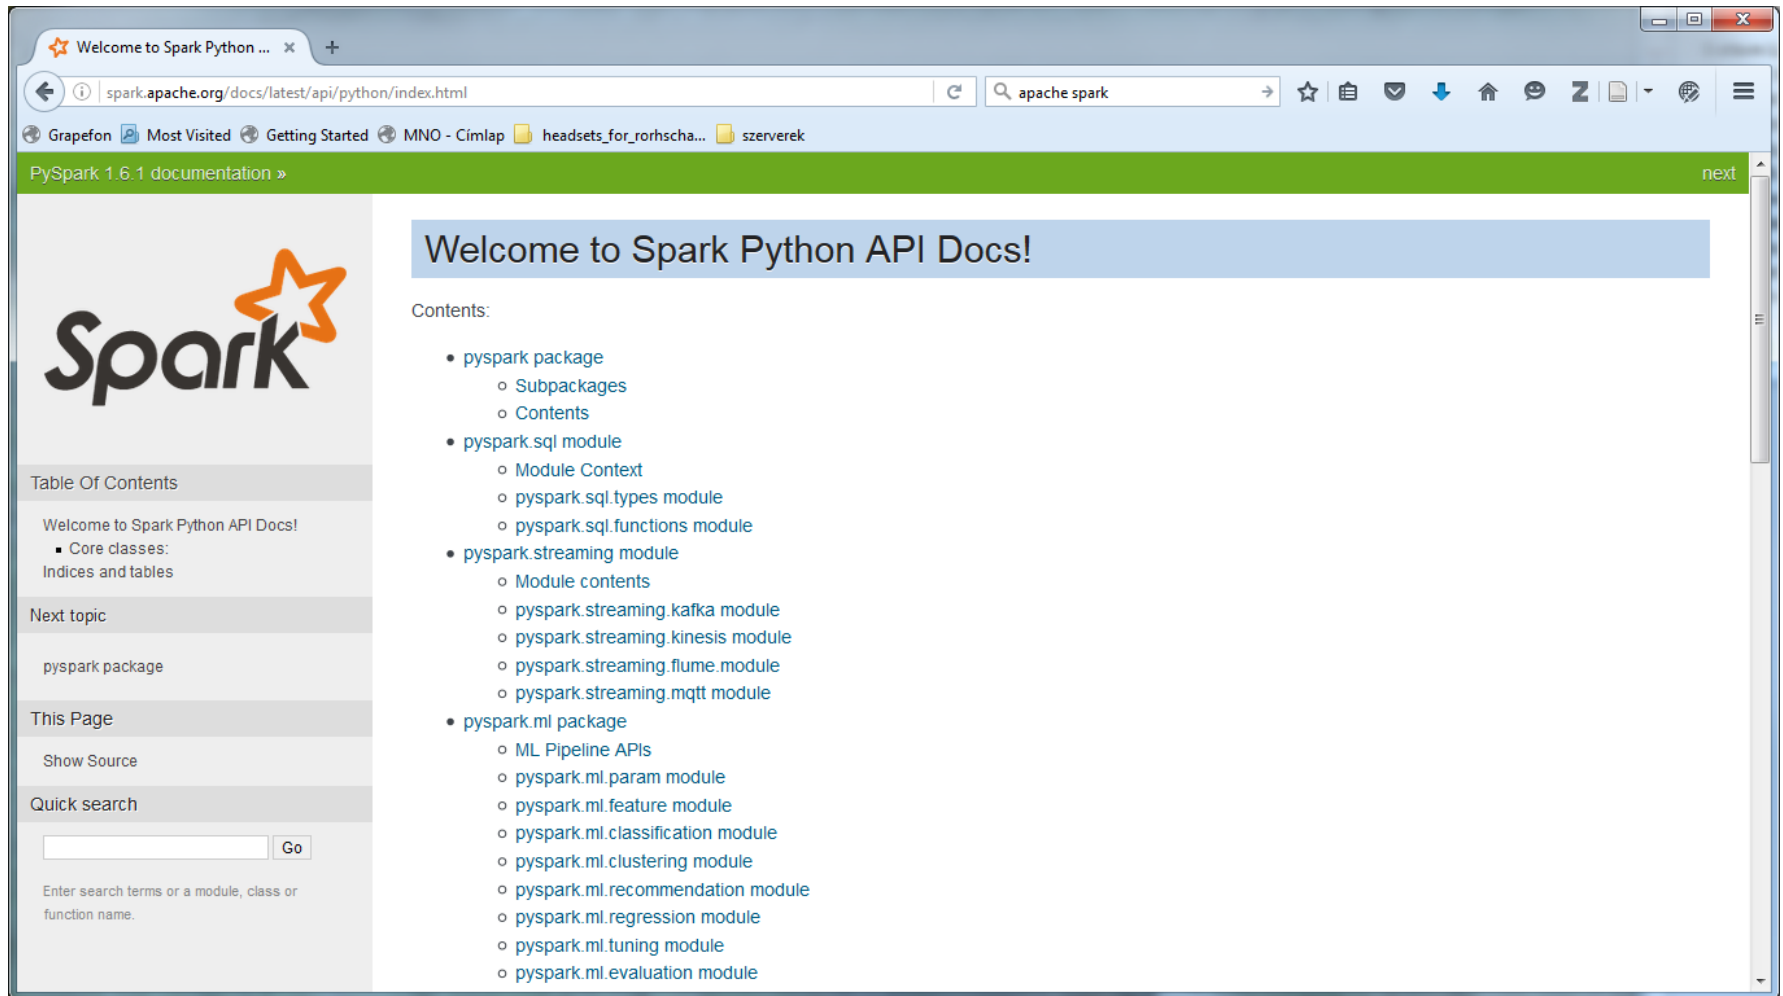
# Spark disztribúciók, cloud

- Hortonworks
- Cludera
- Felhőszolgáltatók

Szolgáltatás	Google	Microsoft	Amazon
SQL adatbázis	+	+	+
NoSQL adatbázis	Hbase	Hbase, DocumentDb	Amazon DynamoDB
Disztribúciók	Hortonworks, Cloudera	Hortonworks beépítve, Cloudera	-
Gépi tanulás	-	+	+
Különböző védelmi eszközök	+	+	+



# Spark - Python



The screenshot shows a web browser window displaying the Spark Python API documentation. The browser's address bar shows the URL `spark.apache.org/docs/latest/api/python/index.html`. The page has a green header bar with the text "PySpark 1.6.1 documentation »" and a "next" link. The main content area has a blue header bar that says "Welcome to Spark Python API Docs!". Below this, there is a "Contents:" section with a bulleted list of links to various parts of the documentation. On the left side, there is a sidebar with a "Table Of Contents" section, a "Next topic" section, a "This Page" section, and a "Quick search" section with a search input field and a "Go" button.

Welcome to Spark Python ... x +

spark.apache.org/docs/latest/api/python/index.html

apache spark

Grapefon Most Visited Getting Started MNO - Cimlap headsets\_for\_rorhscha... szerverek

PySpark 1.6.1 documentation » next

## Welcome to Spark Python API Docs!

Contents:

- [pyspark package](#)
  - [Subpackages](#)
  - [Contents](#)
- [pyspark.sql module](#)
  - [Module Context](#)
  - [pyspark.sql.types module](#)
  - [pyspark.sql.functions module](#)
- [pyspark.streaming module](#)
  - [Module contents](#)
  - [pyspark.streaming.kafka module](#)
  - [pyspark.streaming.kinesis module](#)
  - [pyspark.streaming.flume module](#)
  - [pyspark.streaming.mqtt module](#)
- [pyspark.ml package](#)
  - [ML Pipeline APIs](#)
  - [pyspark.ml.param module](#)
  - [pyspark.ml.feature module](#)
  - [pyspark.ml.classification module](#)
  - [pyspark.ml.clustering module](#)
  - [pyspark.ml.recommendation module](#)
  - [pyspark.ml.regression module](#)
  - [pyspark.ml.tuning module](#)
  - [pyspark.ml.evaluation module](#)

Table Of Contents

Welcome to Spark Python API Docs!

- [Core classes:](#)
- [Indices and tables](#)

Next topic

[pyspark package](#)

This Page


[Show Source](#)

Quick search

[Go](#)

Enter search terms or a module, class or function name.

- Leverage Python & R with Spark



Python & R  
ecosystem  
MPI

#### Bottom Line

10-100X faster performance


- Interact with data in HDFS and Amazon S3 natively from Python
- Distributed computations without the JVM & Python/Java serialization
- Framework for easy, flexible parallelism using directed acyclic graphs (DAGs)
- Interactive, distributed computing with in-memory persistence/caching





# Overview of Dask

**Dask** is a Python parallel computing library that is:

- **Familiar:** Implements parallel NumPy and Pandas objects
  - **Fast:** Optimized for demanding numerical applications
  - **Flexible:** for sophisticated and messy algorithms
  - **Scales up:** Runs resiliently on clusters of 100s of machines
  - **Scales down:** Pragmatic in a single process on a laptop
  - **Interactive:** Responsive and fast for interactive data science
- 

Dask **complements** the rest of Anaconda. It was developed with NumPy, Pandas, and scikit-learn developers.

# Spectrum of Parallelization

Explicit control: Fast but hard

Implicit control: Restrictive but easy



Threads  
Processes  
MPI  
ZeroMQ

Dask

Hadoop  
Spark

SQL:  
Hive  
Pig  
Impala